



Efficient LLM Inference with SGLang

Lianmin Zheng

xAI

Content

SGLang overview

Major techniques

- Efficient KV cache reuse with RadixAttention
- Efficient constrained decoding with compressed finite state machine
- Low-overhead CPU scheduling
- Torch native optimizations (torch.compile, torchao)

Preliminary benchmark results on MI300

Open-source community and roadmap

SGLang Overview

SGLang is a fast serving framework for large language models and vision language models.


What is SGLang?

A **fast inference engine** for LLMs

Comes with its **unique features** for better performance

Serves the **production and research** workloads at xAI



 Ask anything



Grok can make mistakes. Verify its outputs.

Solve the Two
Sum problem in
Python



Recommend a
fantasy RPG
game



Tell me today's
headlines



Help me write a
cover letter



SGLang provides leading inference performance

Compared to the other popular inference engines:

v0.1 (Jan. 2024)

5x higher throughput with automatic KV cache reuse
3x faster grammar-based constrained decoding

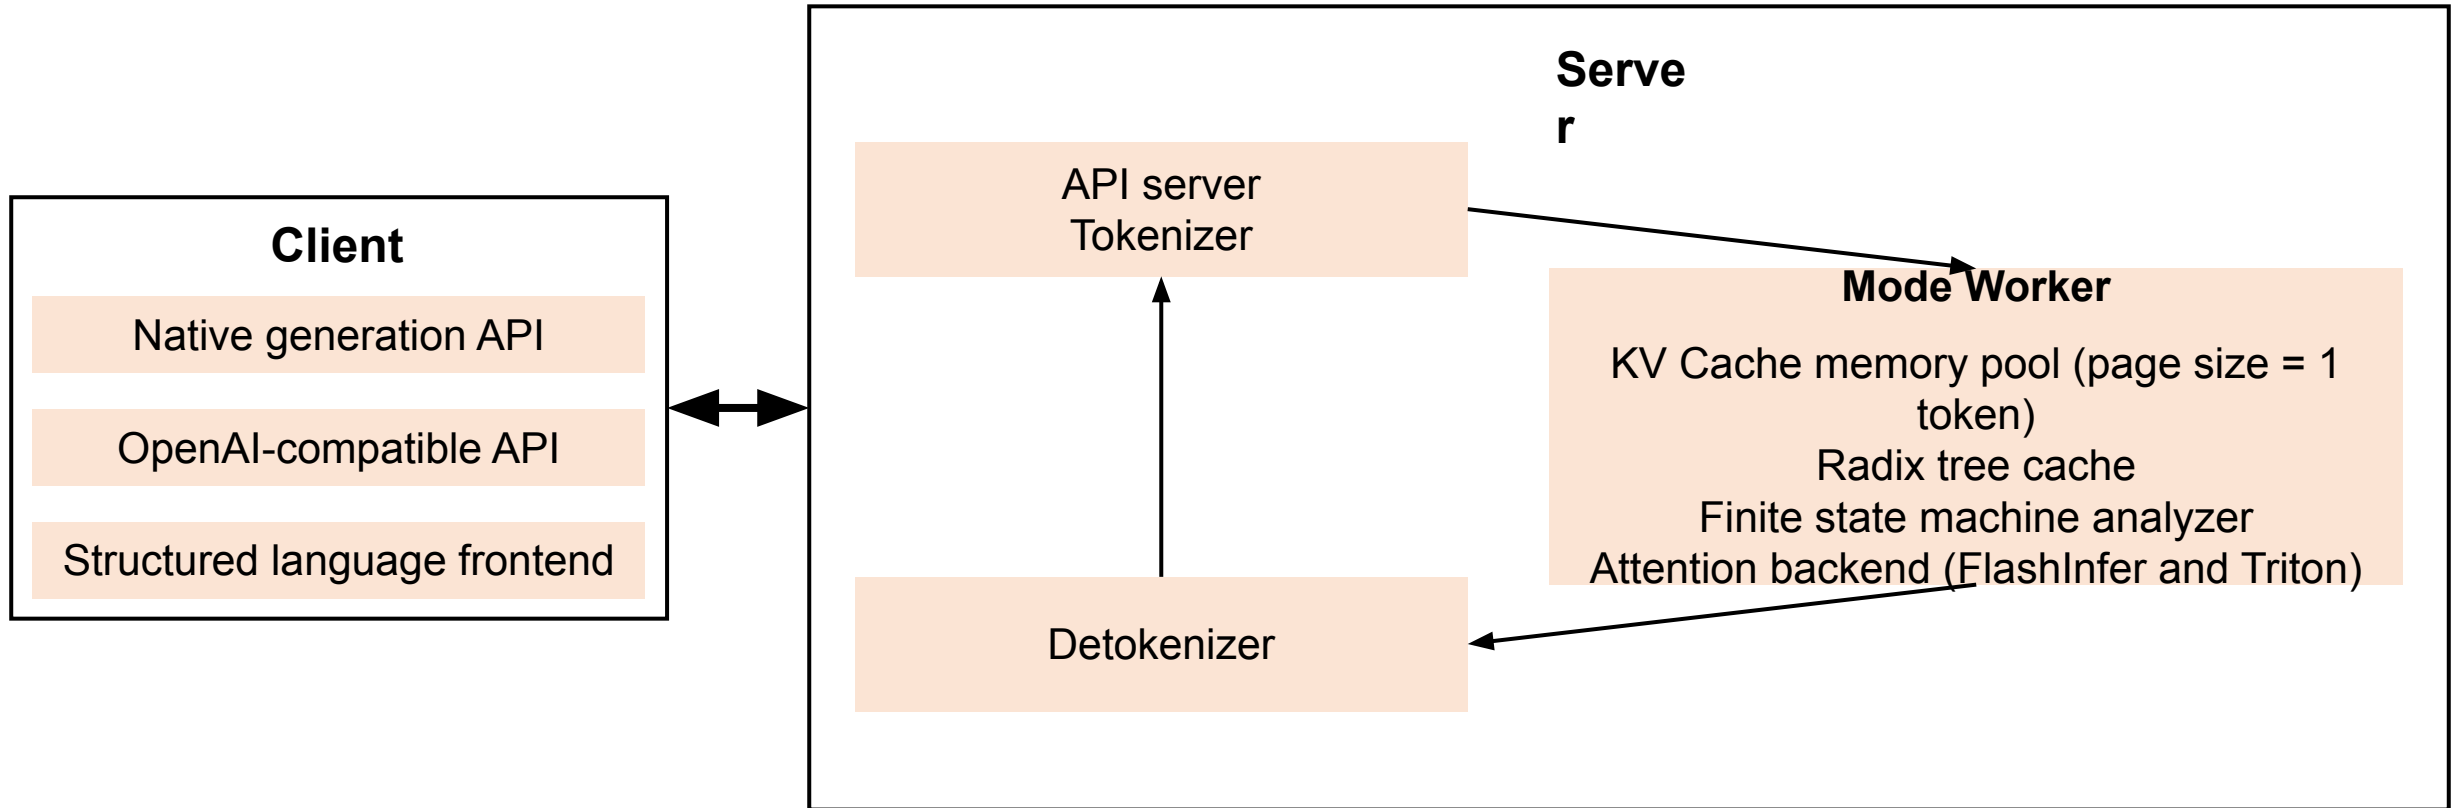
v0.2 (July 2024)

3x higher throughput with low-overhead CPU runtime

v0.3 (Sept. 2024)

7x faster triton attention backend for custom attention variants (MLA)
1.5x lower latency with torch.compile

SGLang architecture overview



Lightweight and customizable code base in
Python/PyTorch

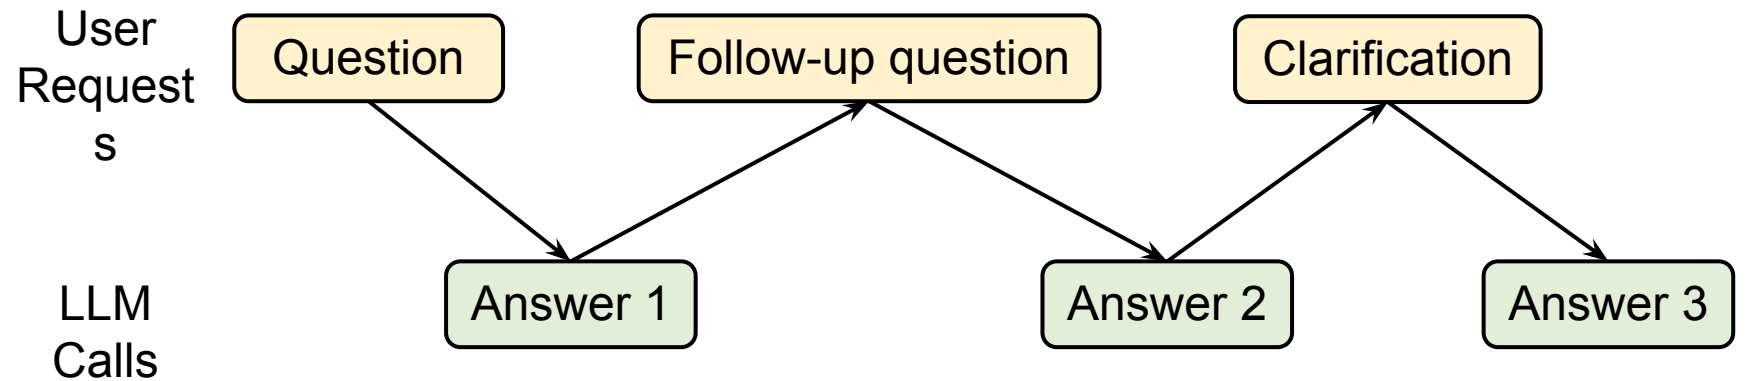
Major Techniques

Four techniques covered in this talk

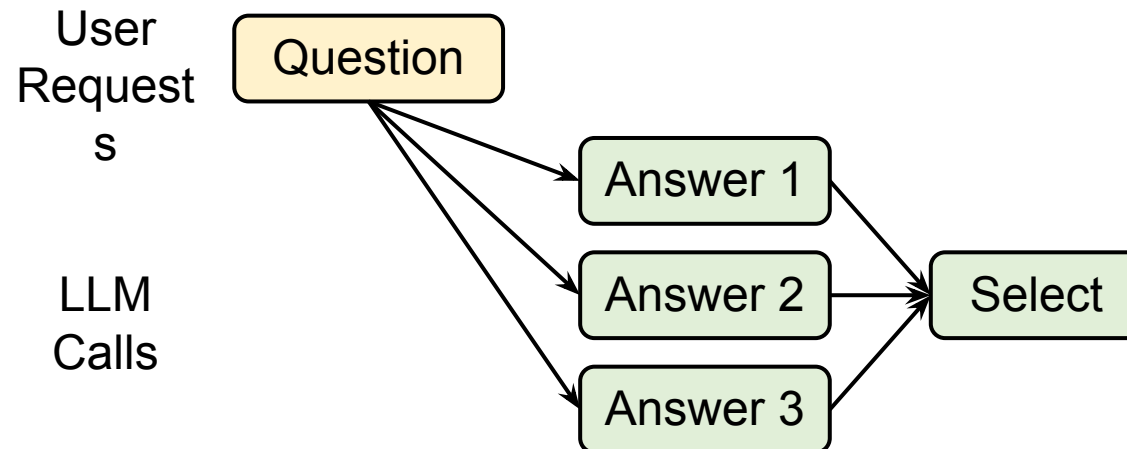
1. Efficient KV cache reuse with RadixAttention
2. Efficient JSON decoding with compressed finite state machine
3. Low-overhead CPU scheduling
4. Torch native optimizations (torch.compile, torchao)

LLM inference pattern: Complex pipeline with multiple LLM calls

Chained calls



Parallel calls



LLM inference pattern: Complex pipeline with multiple LLM calls

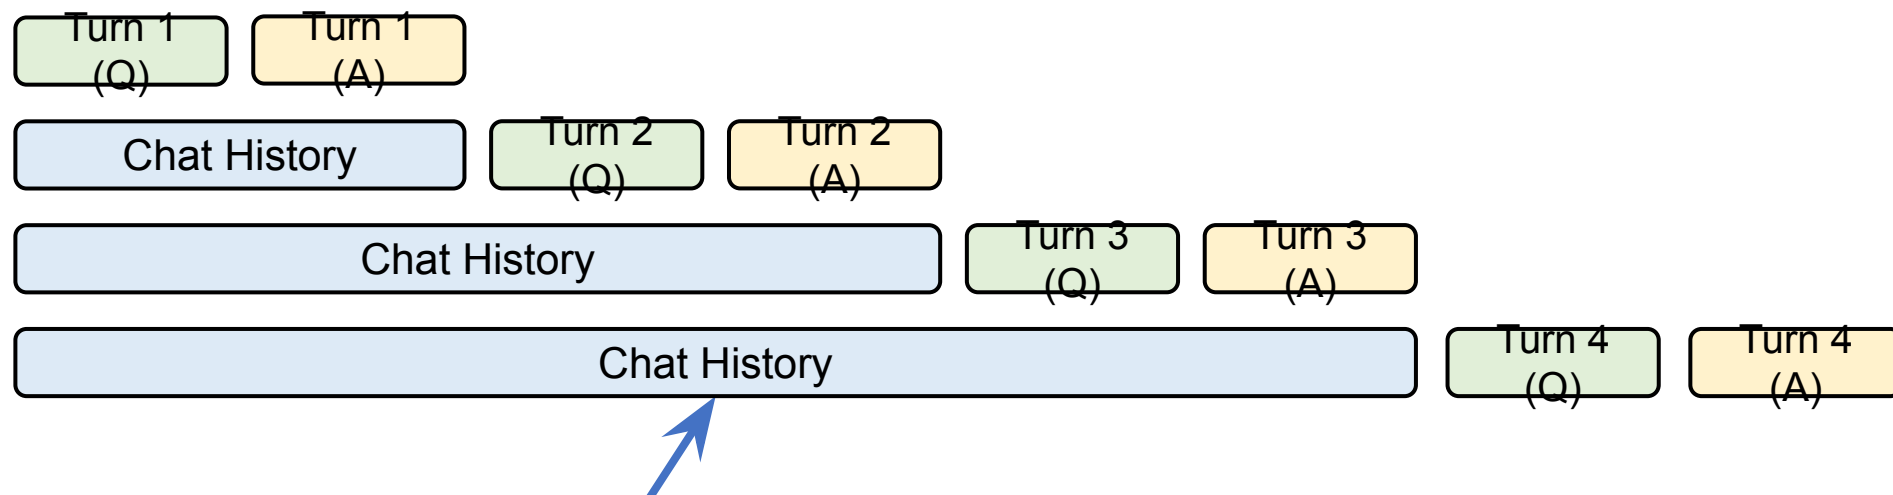
**Chained
calls**

**Parallel
calls**

Multi-call structure **brings optimization opportunities** (e.g., caching, parallelism, shortcut)

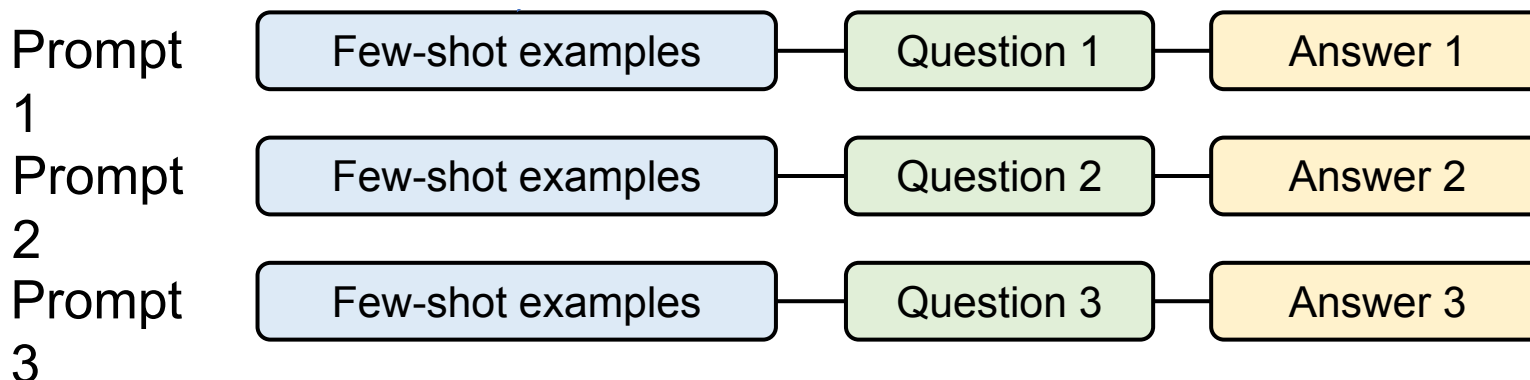
There are rich structures in LLM calls

(a) Multi-turn chat

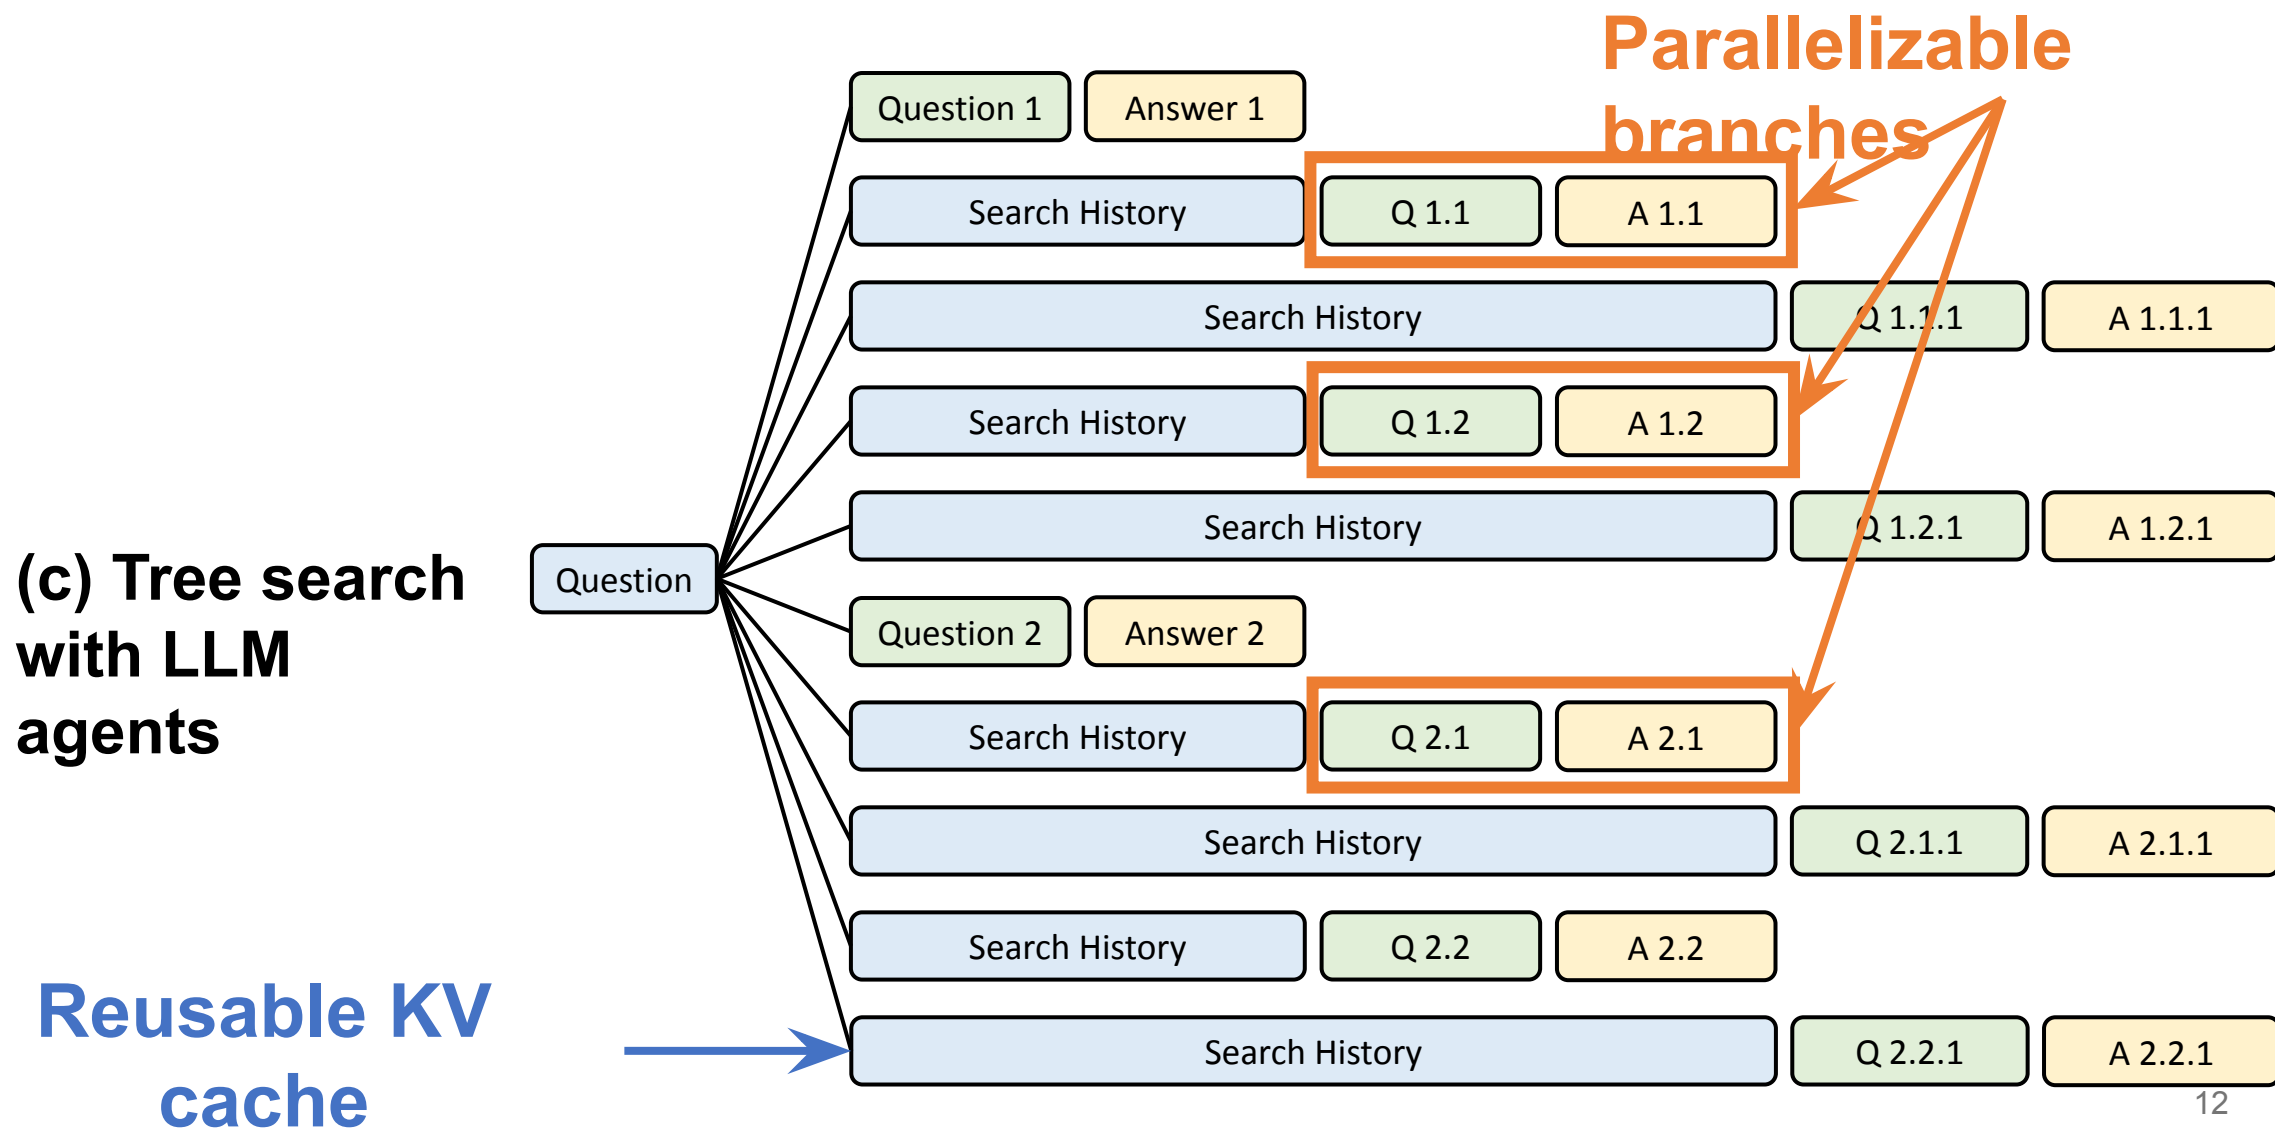


Reusable KV cache (Key-Value cache, some intermediate tensors)


(b)
Few-shot
learning



The structures can be very complicated



Technique 1: Efficient KV cache reuse with RadixAttention



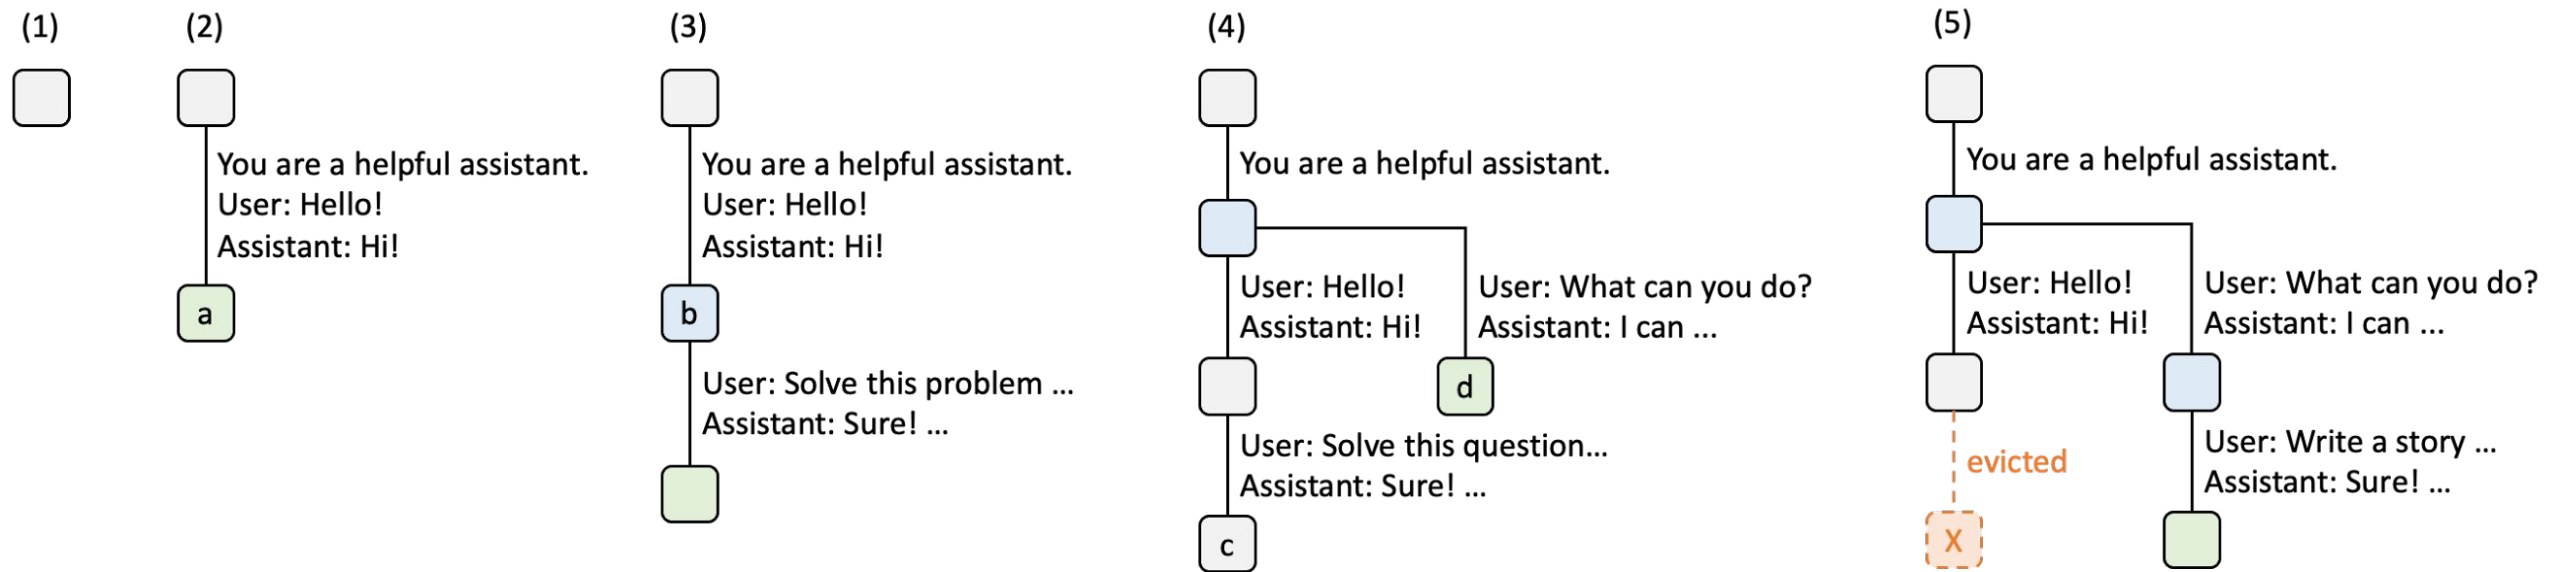
KV Cache

- Some reusable intermediate tensors
- Can be very large (>20GB, larger than model weights)
- Only depends on the prefix tokens

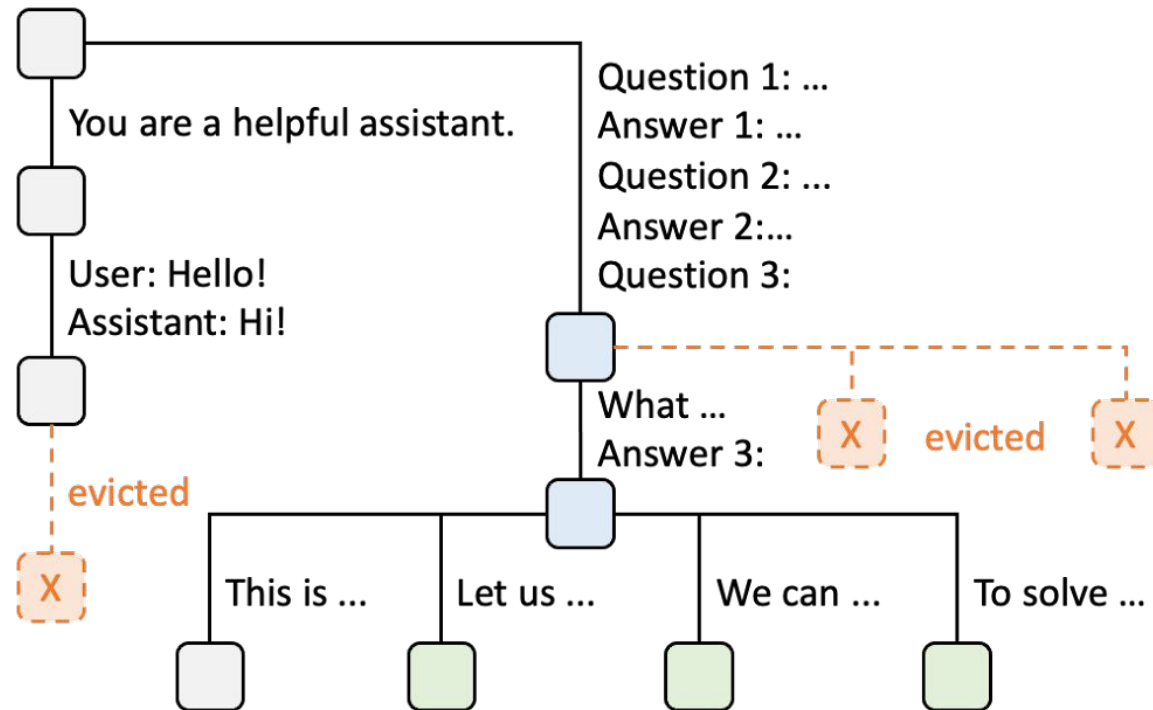
Existing systems: Discard KV cache after an LLM call finishes

Ours: Maintain the KV cache of all LLM calls in a radix tree (compact prefix tree)

RadixAttention maintains the KV cache of all LLM calls in a radix tree (compact prefix tree)



RadixAttention handles complex reuse patterns



RadixAttention enables **efficient prefix matching, insertion, and eviction.**

It handles trees with hundreds of thousands of tokens.

Cache-aware scheduling increases cache hit rate

Idea: Utilize **user annotations** and **runtime metrics** for scheduling

Single worker case


Sort the requests in the queue according to matched prefix length

Distributed case

Route the requests to the worker with the matching cache

Results: SGLang is fast and flexible

- Up to **5x higher throughput** with KV cache reuse and parallelism
- Works automatically across workloads and text/image tokens




Technique 2: Efficient constrained decoding

Workload: Generate the descriptions of characters in the JSON format

SGLang	Outlines + vLLM

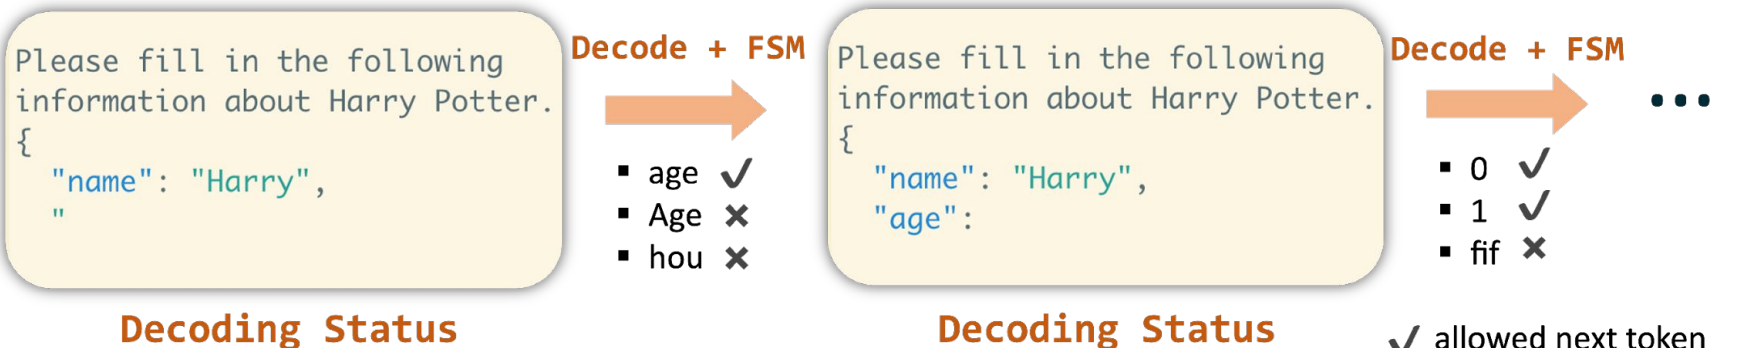
Constrained decoding works by masking the invalid tokens

Constraint decoding: JSON schema -> regular expression -> finite state machine -> logit mask



Regular Expression

Finite State Machine




✓ allowed next token
✗ not allowed next token

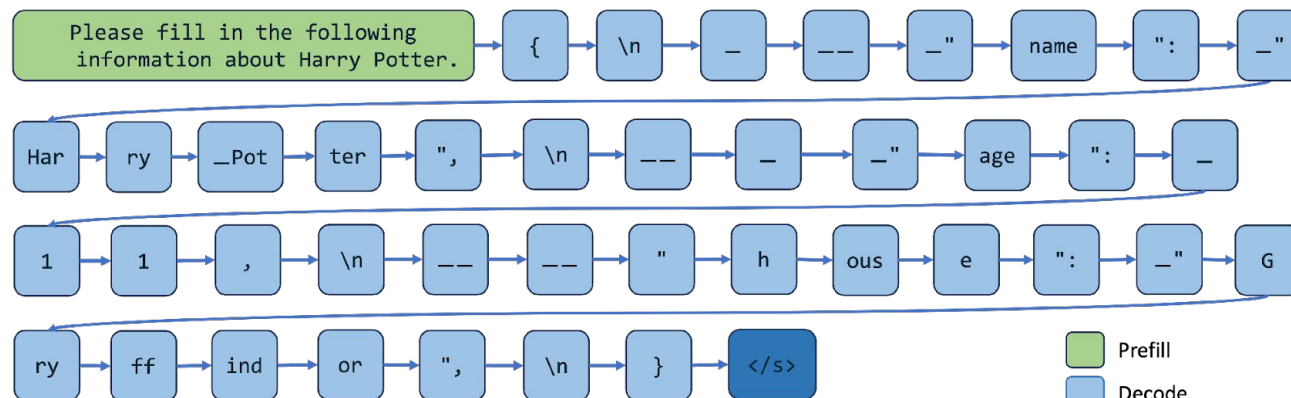
Constrained Decoding With Logits Mask

Compressing the finite state machine allows decoding multiple tokens

We can compress many deterministic paths in the state machine



Jump-Forward Decoding With Compressed FSM



Normal Decoding With FSM




```
Please fill in the following
information about Harry Potter.
{
  "name": "Harry",
  "age": 15,
  "house": "Gryffindor"
}
```

```
Please fill in the following
information about Harry Potter.
{
  "name": "Harry",
  "age": 15,
  "house": "Gryffindor"
}
```

Generated JSONs

Technique 3: Low overhead CPU scheduling

An unoptimized inference engine can waste more than 50% time on CPU scheduling.




Source: https://mlsys.wuklab.io/posts/scheduling_overhead/

Technique 3: Low overhead CPU scheduling

Idea: Vectorize CPU operations / Overlap CPU scheduling

Results: Our python runtime matches C++ runtime and outperforms other python runtime by up to 3x.

Llama-70B (fp8) on 8 GPUs. Higher Throughput is Better.




Source: SGLang v0.2 blog, <https://lmsys.org/blog/2024-07-25-sglang-llama3/>

Technique 4: PyTorch-native optimizations

1.5x faster decoding with torch.compile

1.3x faster decoding with torchao int4 quantization (vs. fp8)




Source: SGLang v0.3 blog, <https://lmsys.org/blog/2024-09-04-sglang-v0-3/>

Preliminary benchmark results on MI300

Grok-1 (314B MoE, FP8) with SGLang on MI300


Setup: both use the triton attention backend. H100 runs TP=8, MI300 runs 2 x TP=4 thanks to its larger memory.
Preliminary results after one week of optimization. MI300 already shows promising results.



Data and integration contributed by the AMD team

Open-source community and roadmap

Community users and contributors



Roadmap

Performance optimizations

Sequence parallelism and sparse attention for long context inference

Adaptive speculative decoding for all batch sizes

Disaggregated prefill and decoding

Hierarchical radix cache

Faster grammar parsing libraries

Communication and CPU overhead overlapping

Modular design

Integrate PyTorch-native optimizations

Community building

Bi-weekly online development meeting

Acknowledgment

AMD technical
support



Xiao
Hai




Anush
Elangovan

Soga Lin, Wun-guo Huang

Core developers: Ying Sheng, Liangsheng Yin, Yineng Zhang, Ke
Rao

Contributors 119

SGLang
open-source
contributors



[+ 105 contributors](#)

Question & Answer

Github: <https://github.com/sql-project/sqlang>

Paper (NeurIPS'24) : <https://arxiv.org/abs/2312.07104>

Welcome to join the [slack](#) and bi-weekly dev meeting!

