

# SGLang v0.2: Faster Interface and Runtime for LLM inference

Aug -  
Dec.  
2023

Early Stage: the “programming LLM” paradigm

Jan. -  
now  
2024

Middle Stage: innovative features and optimizations

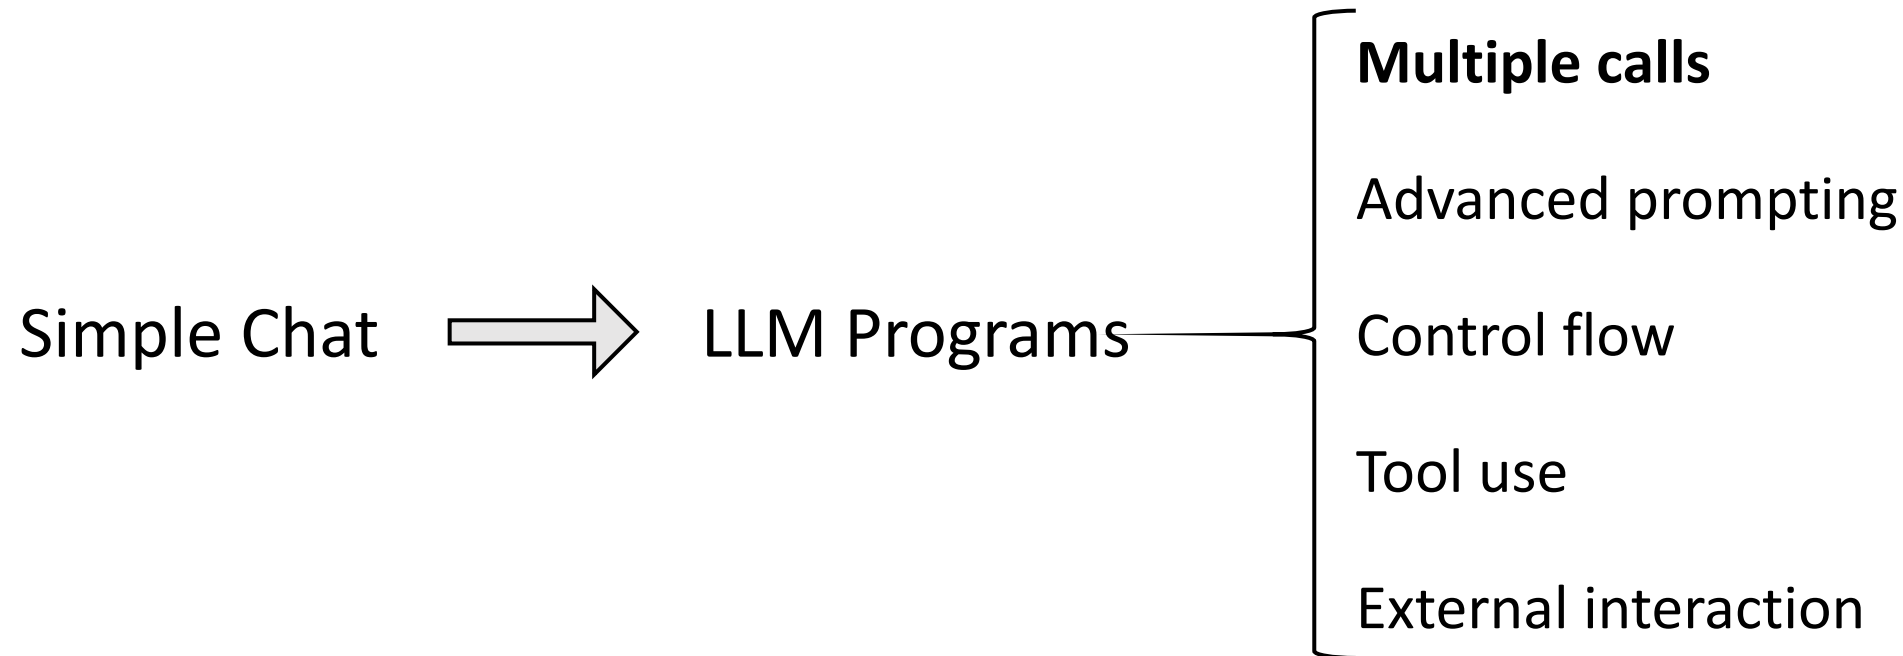
now -  
2024

Production Stage: research and industry use-cases



# Early Stage: the “Programming LLM” Paradigm

From chat and simple prompting to **programmatic usage** of LLMs



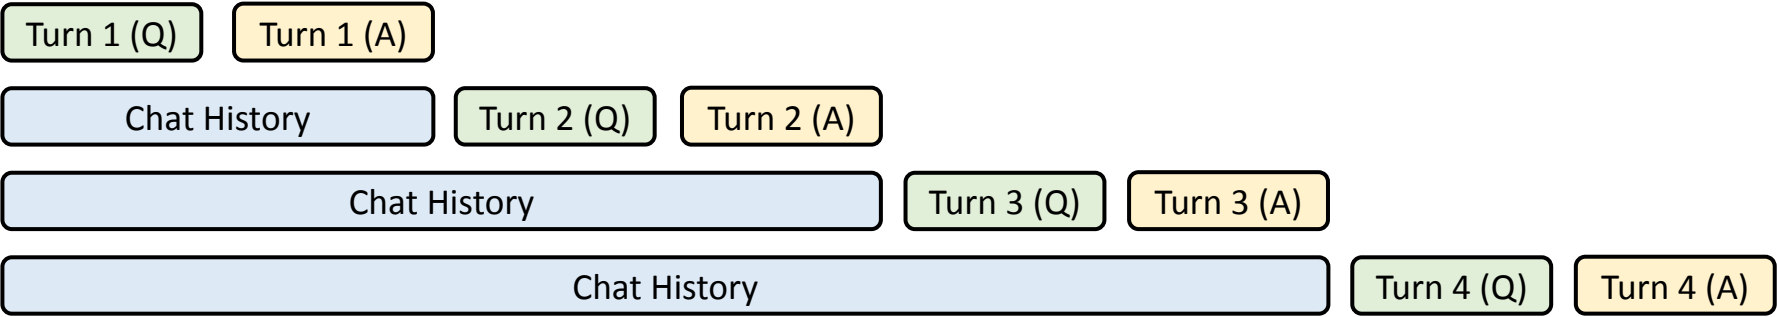
# Existing Systems

**F**ront end language: ignored runtime optimizations  
(Guidance, LMQL)

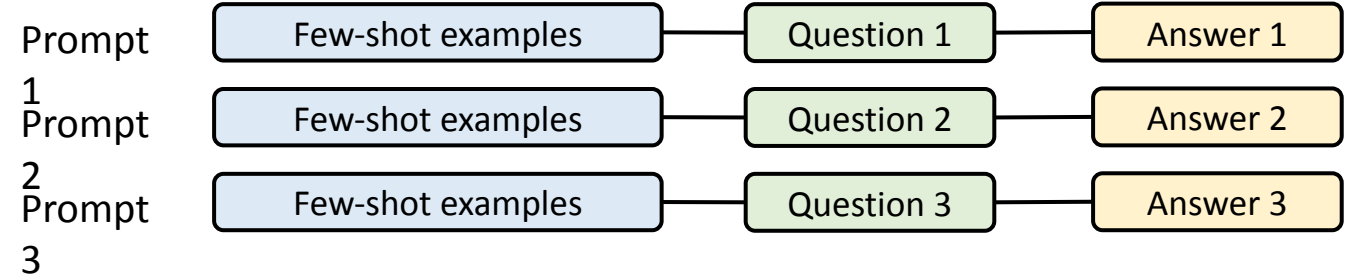
**B**ackend Inference engine: do not know program structure  
(NVIDIA TensorRT-LLM, vLLM)

# Opportunity: KV Cache Reuse


(a) Multi-turn chat



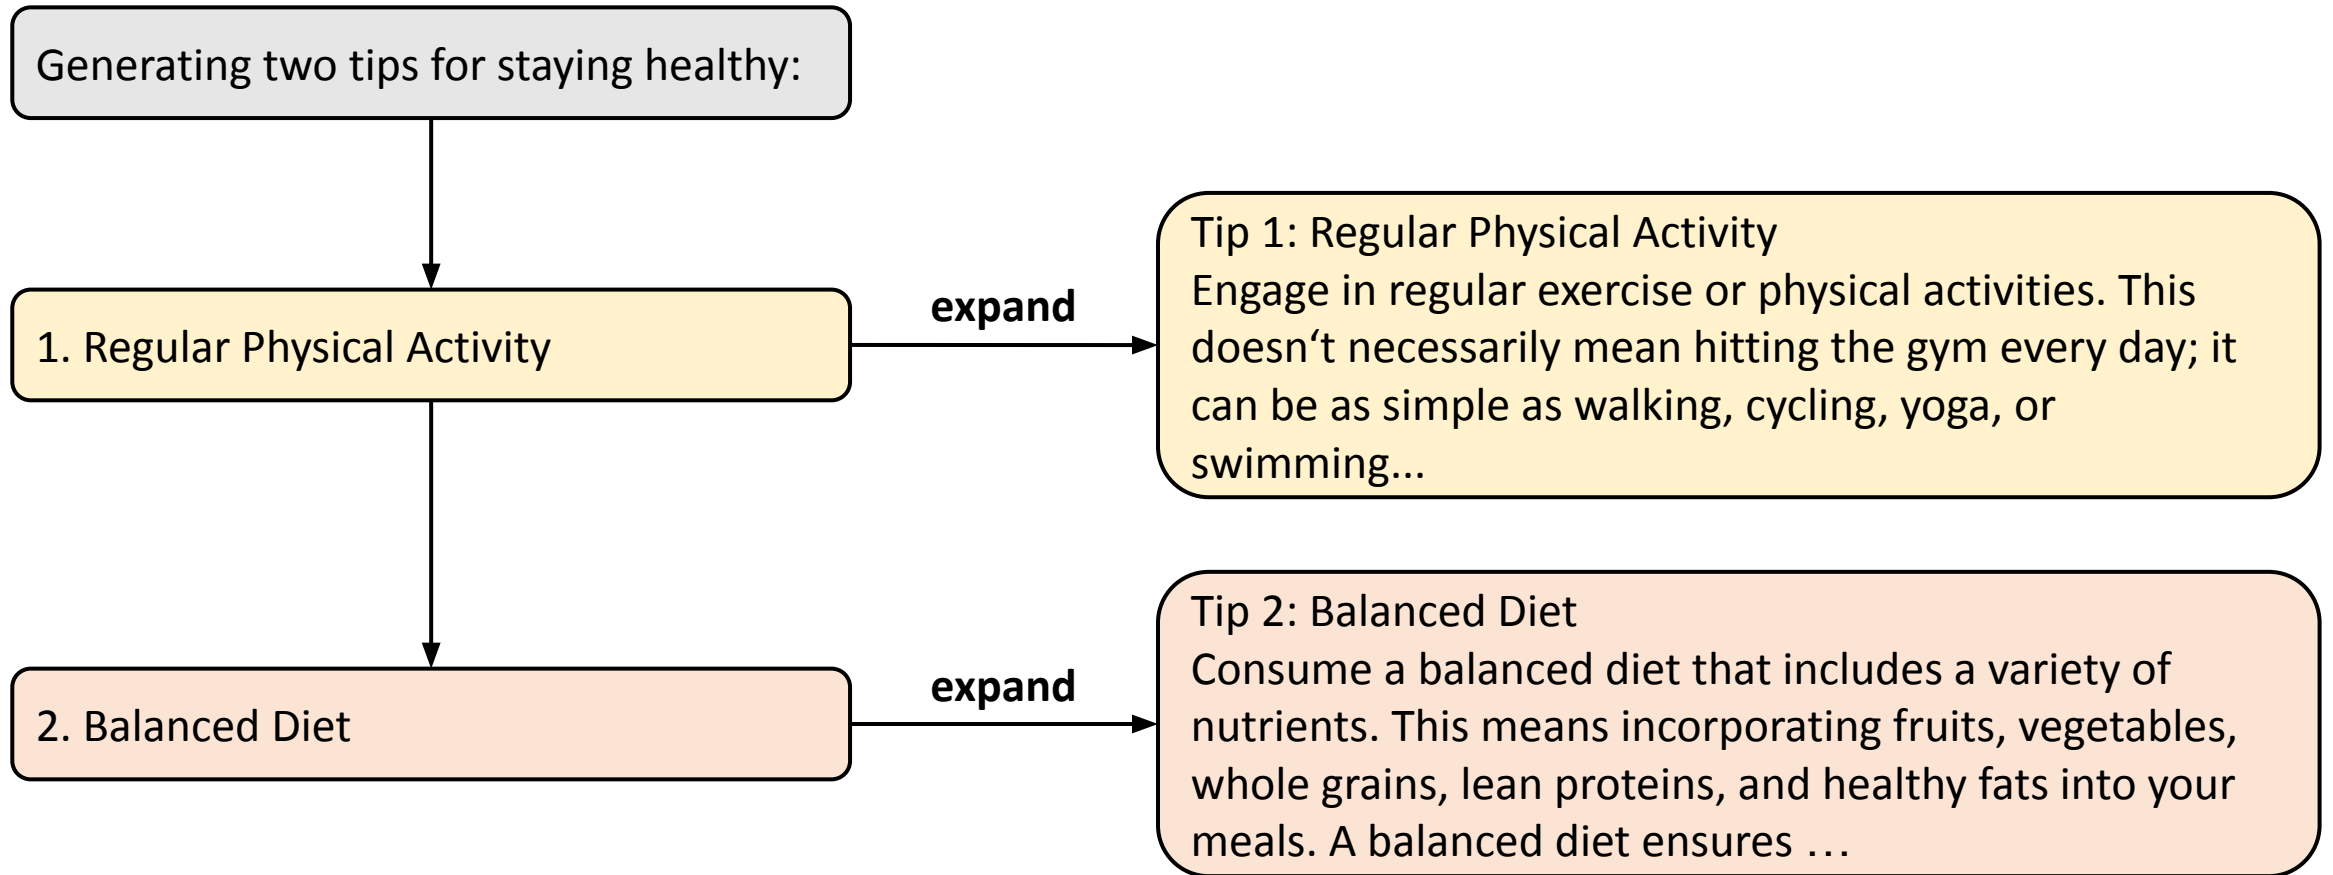
(b) Few-shot learning



(c) Self-consistency



# Opportunity: Parallelism



# System Challenges

- How to program these LLM applications?
- How to optimize across multiple LLM calls?

# Introducing SGLang: A Structured Generation Language

A “co-design” approach

## Front end

- A new domain specific language embedded in Python
- Automatic parallelization and other compiler optimizations

## Back end

- Automatic KV cache reuse with **RadixAttention**



# API example: A Multi-Dimensional Essay Judge

```
dimensions = ["Clarity", "Originality", "Evidence"]
```

```
@function
```

```
def essay_judge(s, essay):
```

```
    s += "Please evaluate the following essay. " + essay
```

```
    # Evaluate an essay from multiple dimensions in parallel
```

```
    forks = s.fork(len(dimensions))
```

```
    for f, dim in zip(forks, dimensions):
```

```
        f += (
```

```
            "Evaluate based on the following metric: " +  
            dim + ". End your judgement with the word 'END'")
```

```
        f += "Judgment: " + f.gen("judgment", stop="END")
```

```
    # Merge judgments
```

```
    for f, dim in zip(forks, dimensions):
```

```
        s += dim + ": " + f["judgment"]
```

```
    # Generate a summary and give a score
```

```
    s += "In summary," + s.gen("summary")
```

```
    s += "I give the essay a letter grade of " +
```

```
    s += s.gen("grade", choices=["A", "B", "C", "D"])
```

```
ret = essay_judge.run(essay="A long essay ...")
```

```
print(ret["grade"])
```

Frontend

Launch parallel prompts

Non-blocking generation call

Fetching generation results

Constrained generation

Run the function

# Compiler Optimizations

- **Building a dataflow graph**

- Remove Python Interpreter Overhead
- Global scheduling optimization over the graph

Frontend



Backend

- **Prefetching cached prefixes**

- Insert prefetching nodes into the graph

- **Code movement for increasing sharable prefix length**

- Reorder some prompt elements with the help of GPT-4

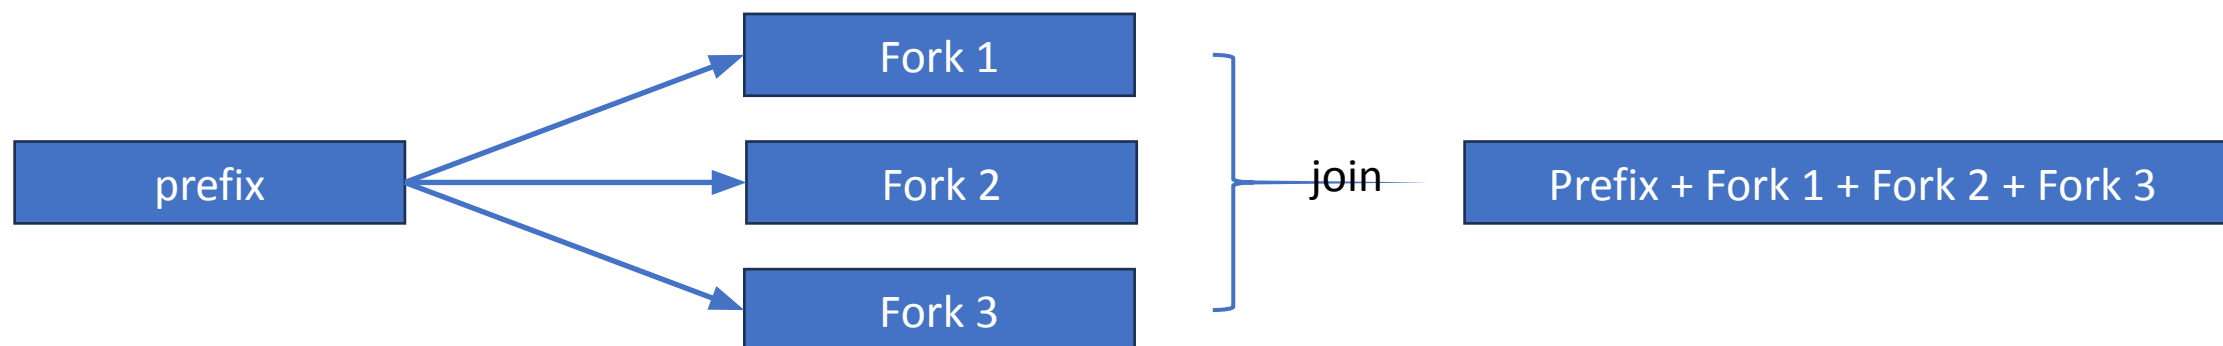
# Prefix caching from request tracking?

- In multi-turn chat, retrieval tasks, etc
  - The interpreter tracks the request id (rid) and caches the history before it ends.
  - Only needs to match the rid.
  - “pin” is a primitive of fixing a prefix to be cached.
  - “fork/join” primitives

Frontend



Backend



Cannot reuse shared prefix across requests!

Aug -  
Dec.  
2023

Early Stage: the “programming LLM” paradigm

Jan. -  
now  
2024

Middle Stage: innovative features and optimizations

Focused efforts on backend/runtime performance

now -  
2024

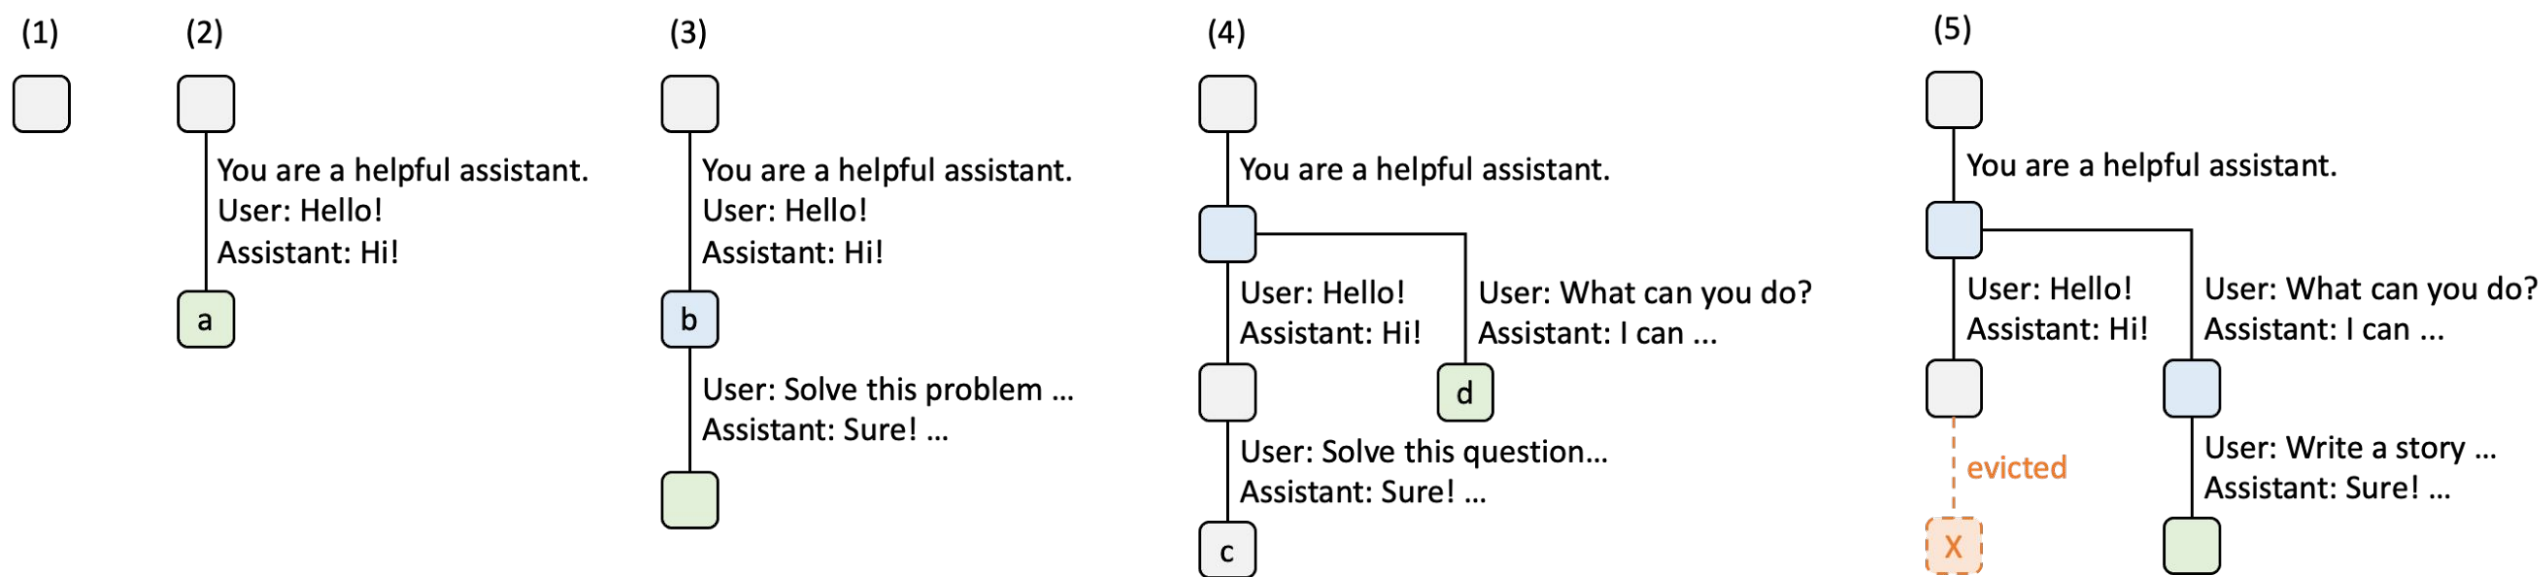
Production Stage: research and industry use-cases



# Runtime (SRT) with RadixAttention

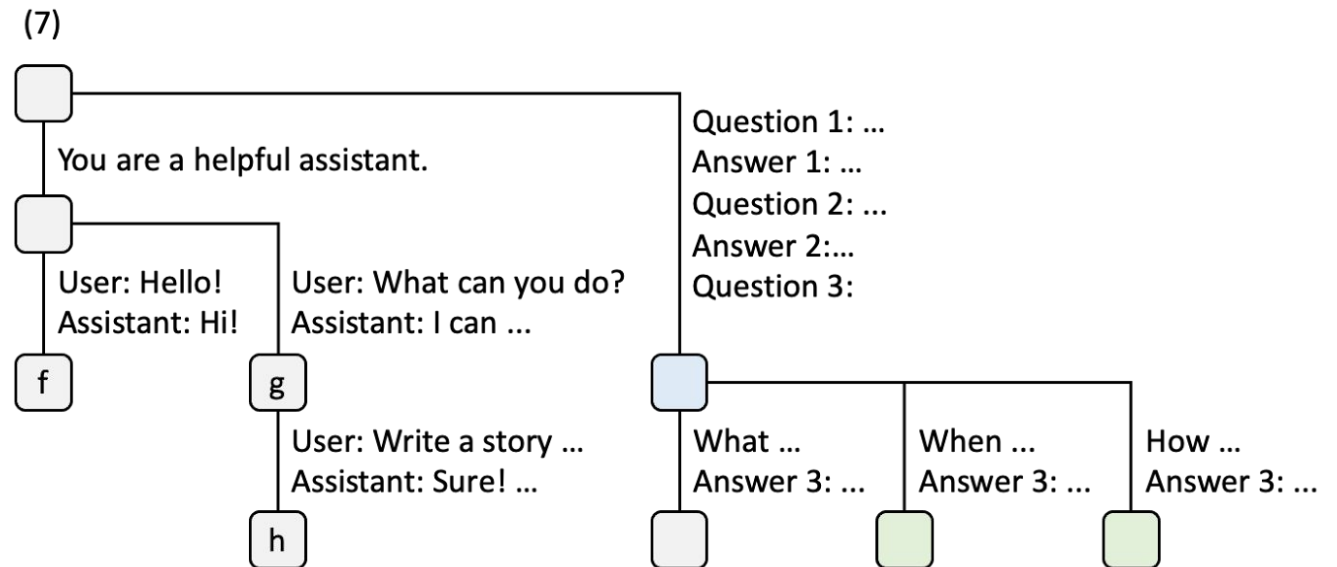
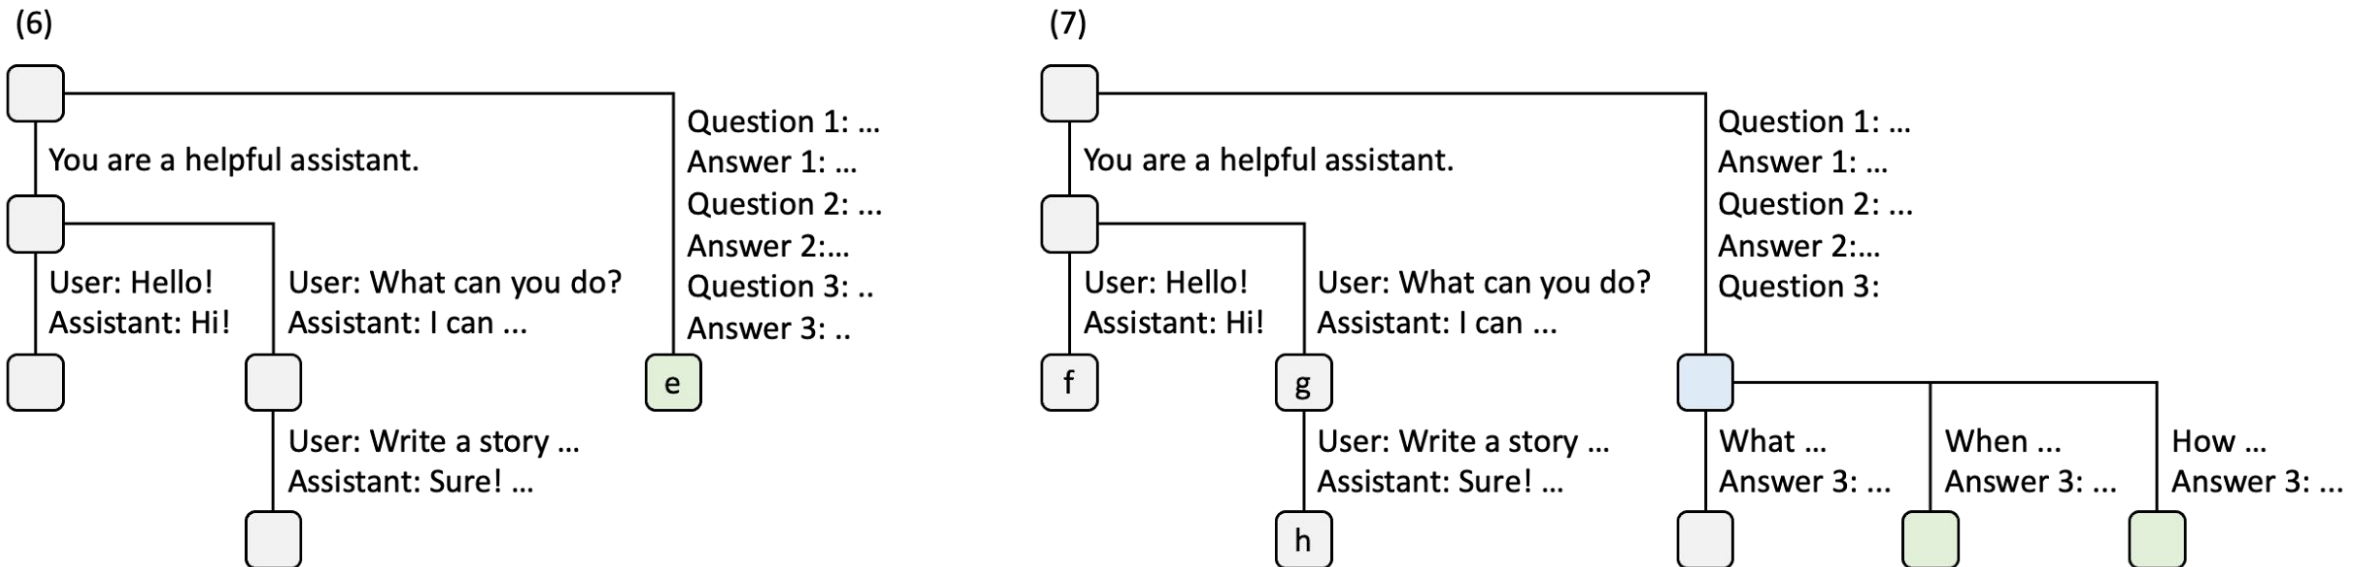
**Existing Systems:** Discard KV cache after a request finishes.

**Ours:** Maintain an LRU cache of the KV cache of all requests in a radix tree (compact prefix tree).



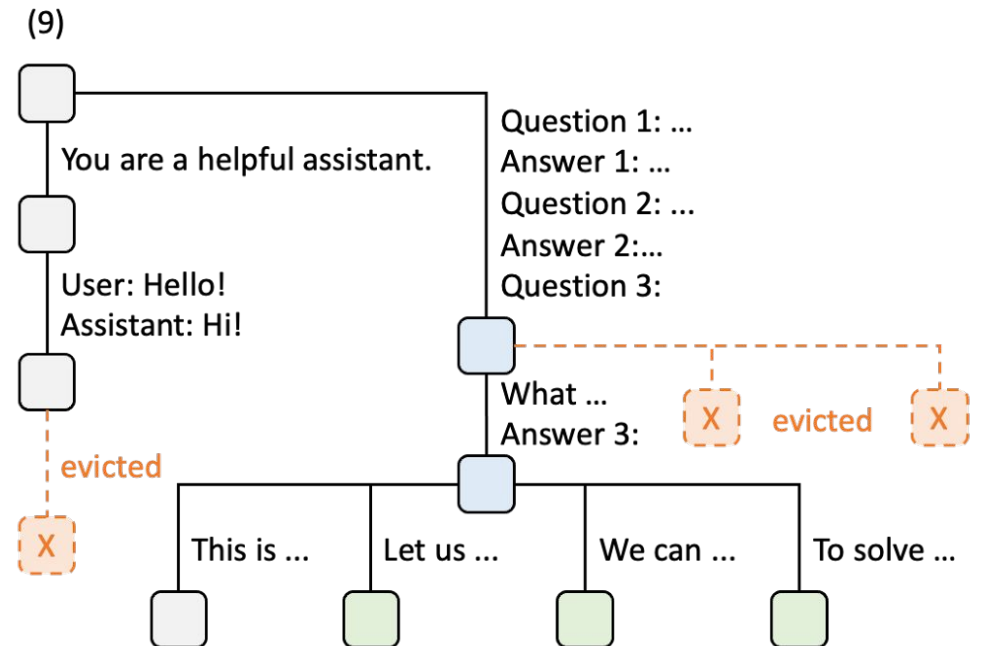
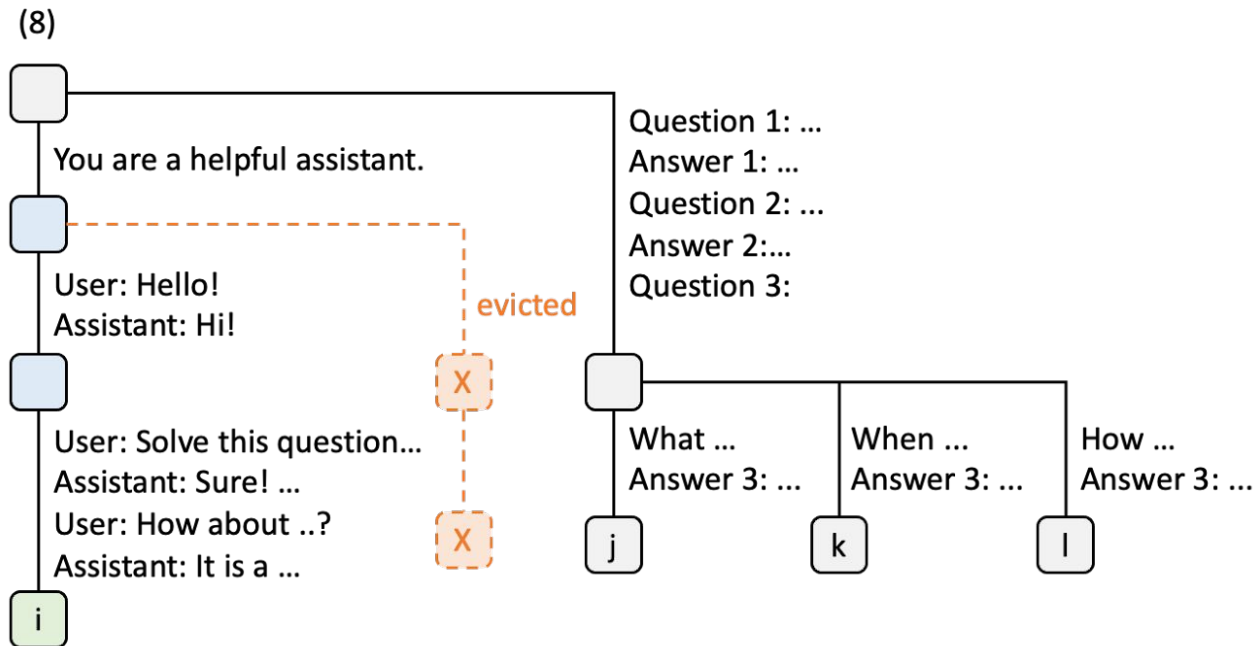
# Runtime (SRT) with RadixAttention

Maintain an LRU cache of the KV cache of all requests in a radix tree.



# Runtime (SRT) with RadixAttention

Maintain an LRU cache of the KV cache of all requests in a radix tree.



# Cache Aware Scheduling

- In the request queue, sort the requests according to the matched prefix length
  - Achieves good cache hit rate
- Future work
  - Distributed cache aware scheduling for multiple data parallel workers
  - Fairness to prevent starvation (<https://arxiv.org/abs/2401.00588>)



Aug -  
Dec.  
2023

Early Stage: the “programming LLM” paradigm

Jan. -  
now  
2024

Middle Stage: innovative features and optimizations

RadixAttention

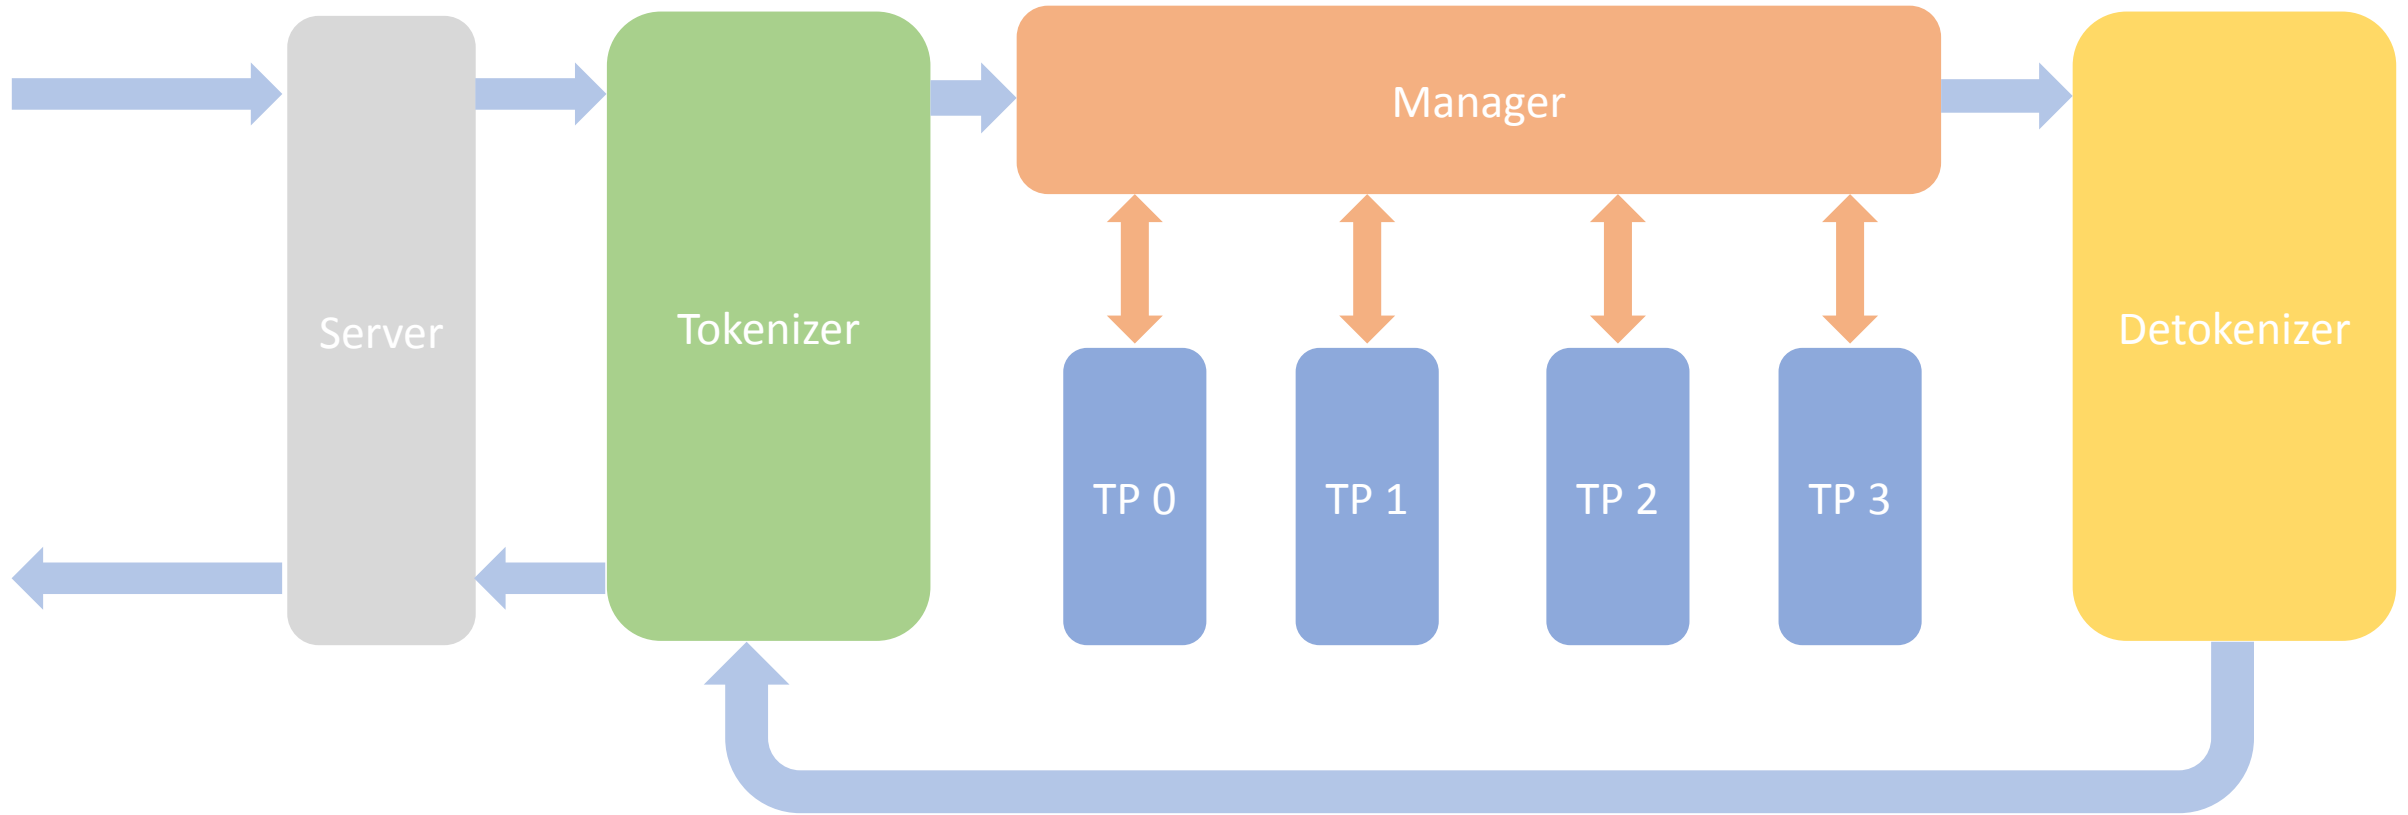
Upper-level Scheduling

now -  
2024

Production Stage: research and industry use-cases




# SGLang Structure: Pipeline



# SGLang Structure: Inside TP Worker

1.Decode Batch




EOS token, finished this req

2.Prefill a New Batch



Merge into decode batch  
ready for decode


3.Decode Batch



- newly decoded token
- prompt tokens
- EOS (newly decoded)
- decoded token

How to always keep the batch size large enough?

# Dynamically Adjust the new token ratio estimation

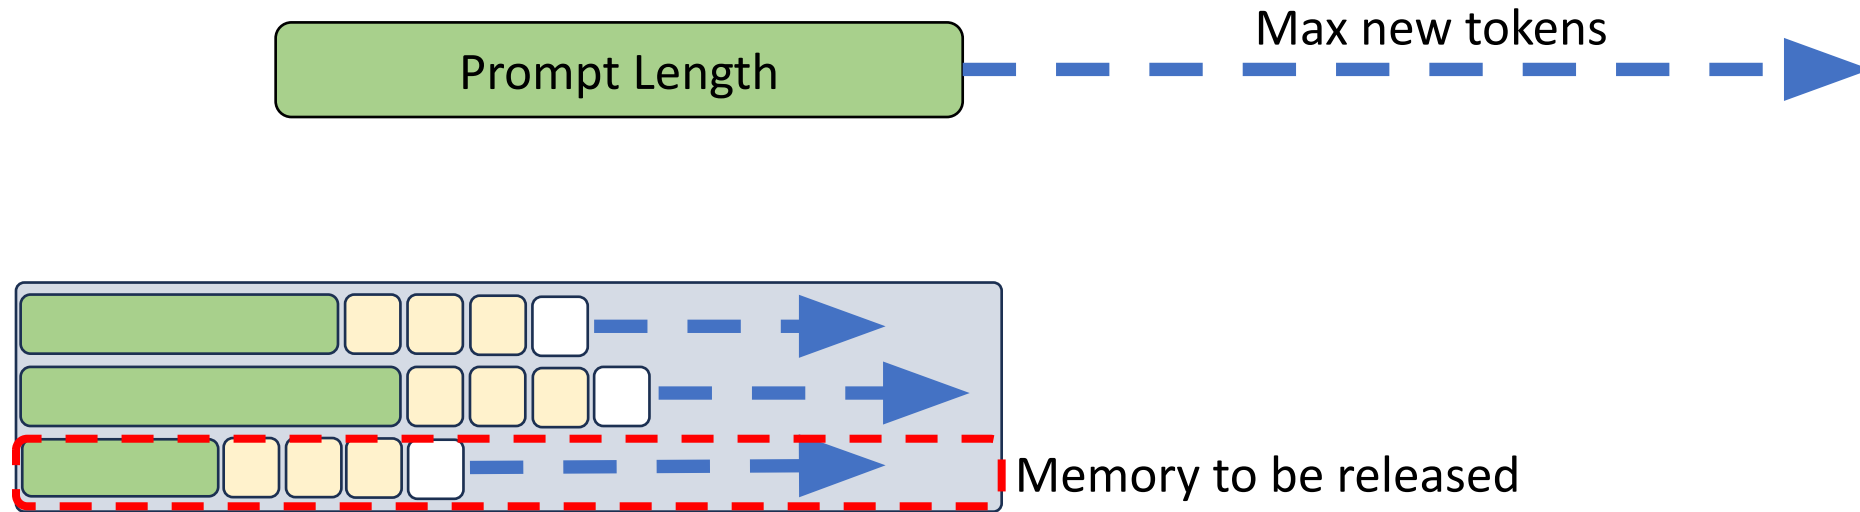


The max context length  
decided by max new tokens



- There is a lot of space left in the GPU memory
- We do not need to reserve every token in max new tokens

# Dynamically Adjust the new token ratio estimation



1. The EOS would be earlier than the max new tokens.
2. There are always requests finished and release all the memory.

Only preserve  $\beta \times \text{max\_new\_token}$  tokens in advance, and adjust  $\beta$  dynamically.

Aug -  
Dec.  
2023

Early Stage: the “programming LLM” paradigm

Jan. -  
now  
2024

Middle Stage: innovative features and optimizations

RadixAttention

Upper-level Scheduling

Jump-forward decoding

now -  
2024

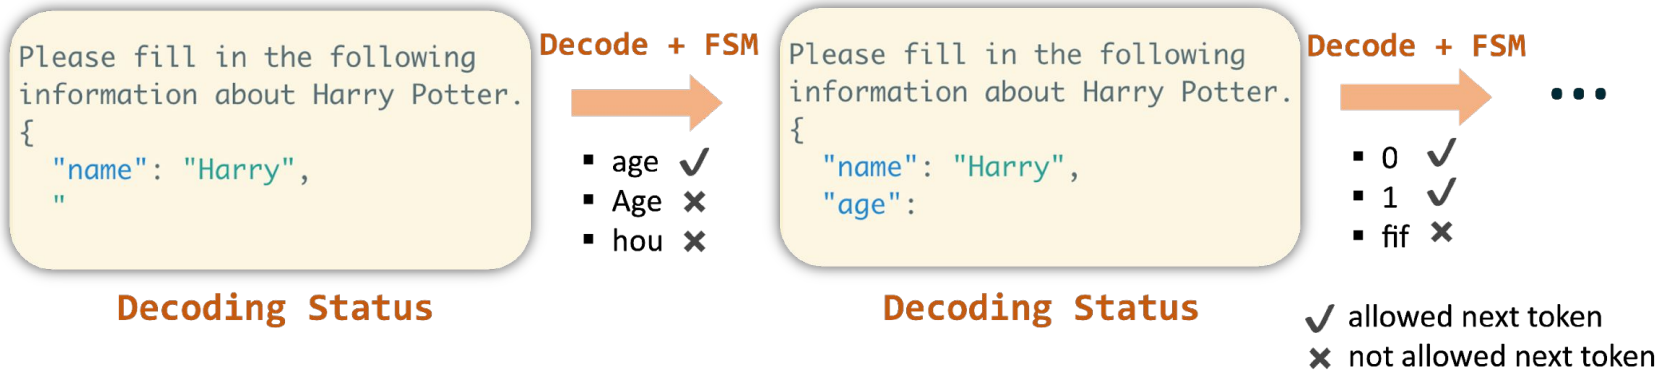

Production Stage: research and industry use-cases




# Jump-forward JSON Decoding

## Method

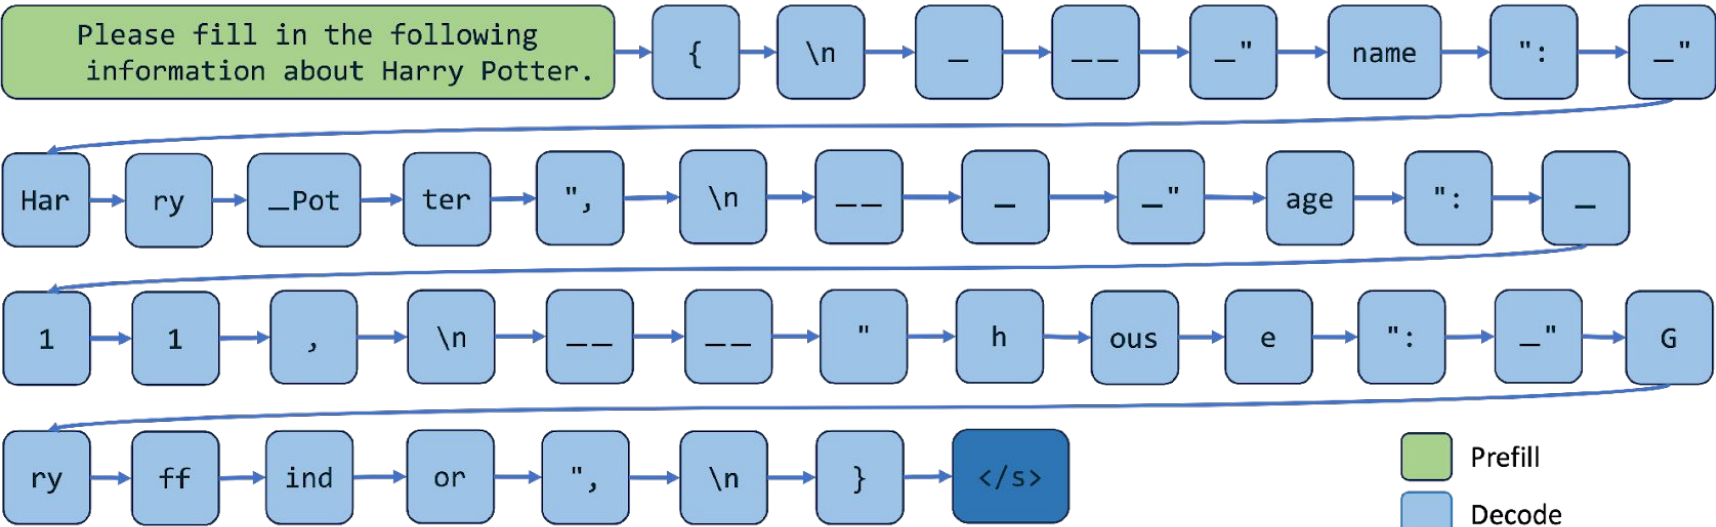
- Analyze the regular expression
- Compress the finite state machine
- Decode multiple tokens at the same time



# Speedup Regex Guided Generation



Jump-Forward Decoding With Compressed FSM



Normal Decoding With FSM

- Prefill
- Decode
- Jump-Forward

```
Please fill in the following
information about Harry Potter.
{
  "name": "Harry",
  "age": 15,
  "house": "Gryffindor"
}
```

```
Please fill in the following
information about Harry Potter.
{
  "name": "Harry",
  "age": 15,
  "house": "Gryffindor"
}
```

Generated JSONs




# Jump-forward JSON Decoding

## Results:

3x faster latency

2.5x higher throughput




# Summary: techniques in SGLang


- RadixAttention
- Jump-forward JSON Decoding
- Torch Compile
- Flashinfer Kernels
- Chunked Prefill
- Continuous Batching
- Token Attention(Paged Attention with `page_size = 1`)
- CUDA Graph
- Interleave window attention

# SGLang v0.2 Results

Llama-8B (bf16) on 1 x A100. Higher Throughput is Better.



Llama-70B (fp8) on 8 x H100. Higher Throughput is Better.



Aug -  
Dec.  
2023

Early Stage: the “programming LLM” paradigm

Jan. -  
now  
2024

Middle Stage: innovative features and optimizations

now -  
2024

Production Stage: research and industry use-cases


# Research and industry use cases



x.ai: Production serving of grok-2 and grok-2-mini on X



Databricks: accelerate research workflow by 3x



LMSys Chatbot Arena: serving vision language models

LLaVA OneVision: serving multi-modal image and video models



# Future work

- multi-level cache
- distributed radix attention
- long-context
- speculative decoding
- communication overlapping
- .....

Do the serving engines come to converge on performance?

YES and NO

Basic performance eventually converge

But there are more sophisticated workloads from different scenarios:  
RAG systems, agent systems, ...

We never forget about the origin of SGLang!  
Structured inputs, interactions with different resources, multi-modality, ...

# Principles in future development

Simplism

Minimalism

Modularity

Ease of use

Development velocity

Performance